

# Head Pose Estimation with MediaPipe and OpenCV in Javascript



Susanne Thierfelder · [Follow](#)

5 min read · Nov 15, 2022

90

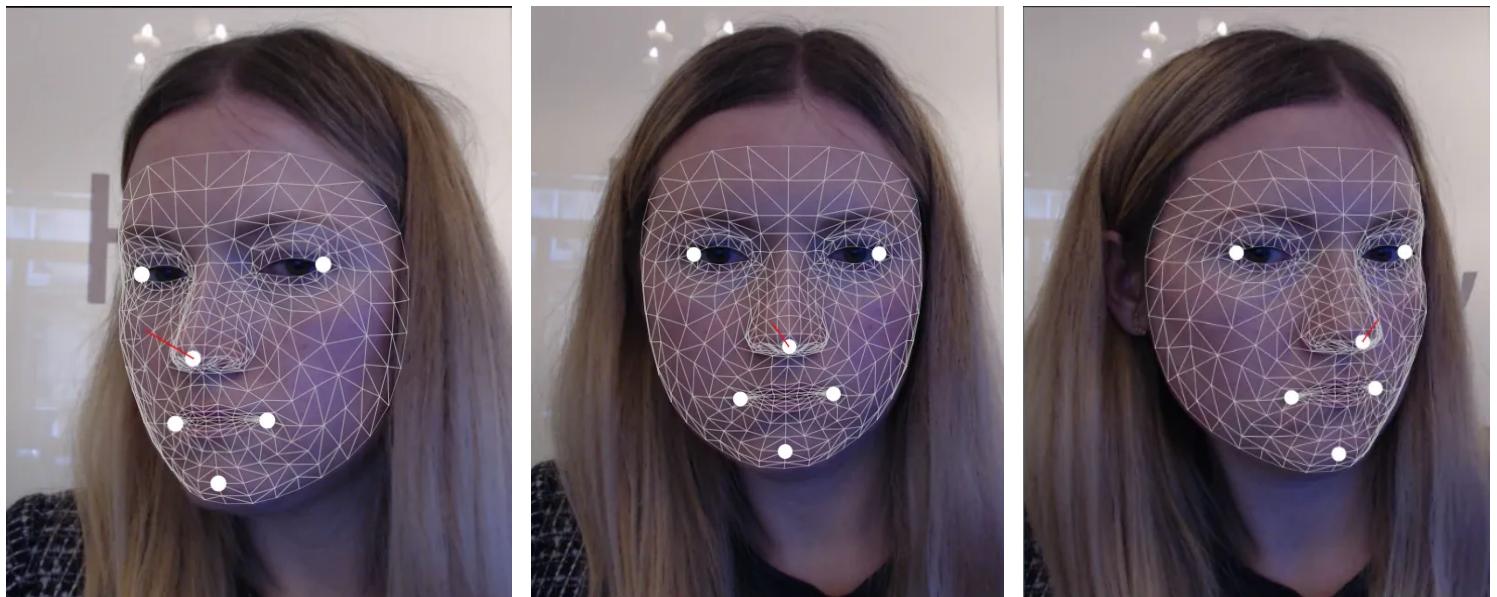
1

+

▶

↑

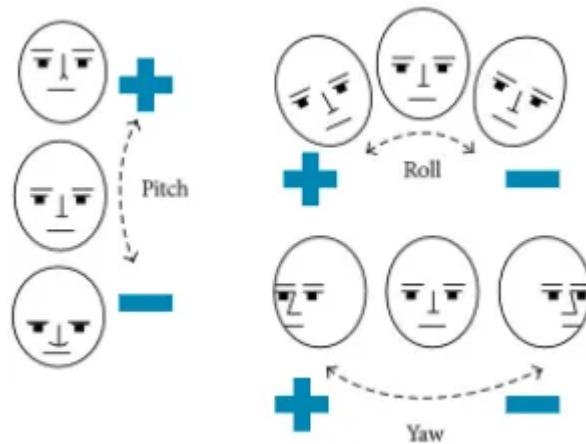
In this blog post, I demonstrate how to estimate the head pose from a single image using MediaPipe FaceMesh and OpenCV in Javascript. Check out [my demo on CodePen!](#)



## Landmarks for pose estimation

Currently, MediaPipe [JavaScript Solution API](#) does not include the option ‘enableFaceGeometry’ which allows obtaining the pose of each detected face automatically (see issue #[2673](#)).

In this article, I will look at the principles of pose computation from 3D-2D point correspondence, explain the relevant OpenCV algorithms, talk about landmarks in MediaPipe, and finally show you how to display the pose of the face.

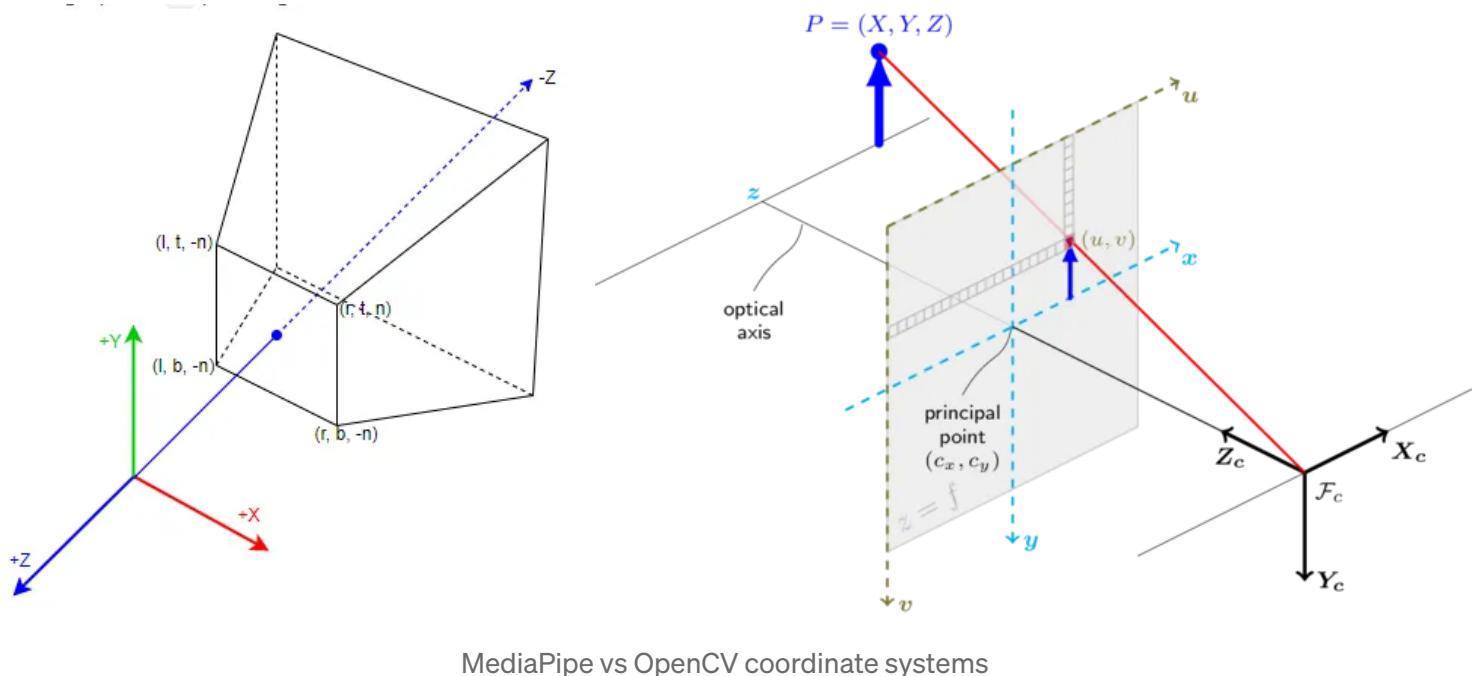


Roll, pitch, and yaw angles for head pose estimation

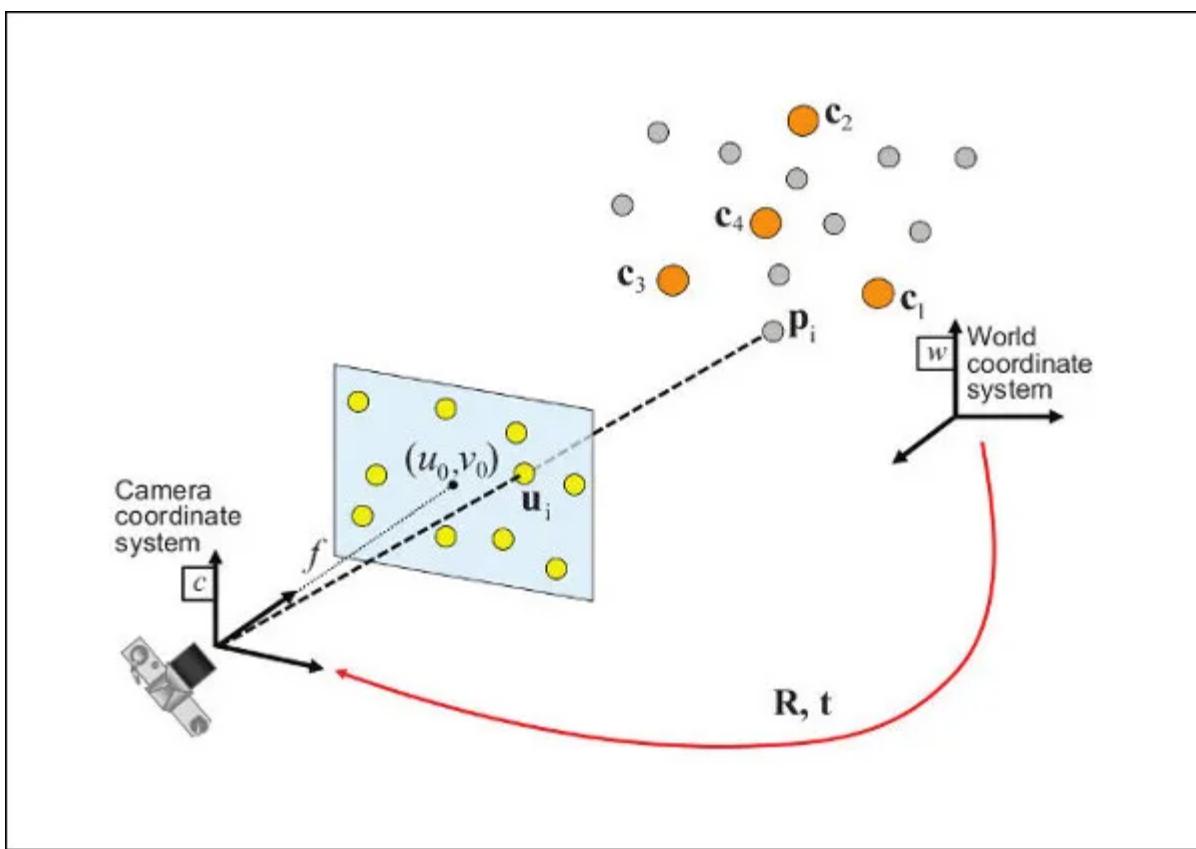
Steps to estimate the face's yaw, pitch, and roll angles in a given image :

- 1) Find face landmarks using Mediapipe ‘FaceMesh’
- 2) Produce rotation vector with OpenCV-Javascript *solvePnP* function
- 3) Pass rotation vector to OpenCV *Rodrigues* function to get rotation matrix
- 4) Finally, decompose the rotation matrix to get Euler angles

The **pose computation problem** consists in solving for the rotation and translation that **minimizes the reprojection error from 3D-2D point correspondences** ([OpenCV](#), see also [publication](#)). The reprojection error is a geometric error corresponding to the image distance between a projected point and a measured one (see [Wikipedia](#)).



`cv.solvePnP`: Finds an object pose from 3D-2D point correspondences. This function returns the rotation R and the translation vectors t that transform a 3D point expressed in the world frame into the camera frame:

Pose computation overview ([OpenCV](#))

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \mathbf{A} \boldsymbol{\Pi}^c \mathbf{T}_w \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

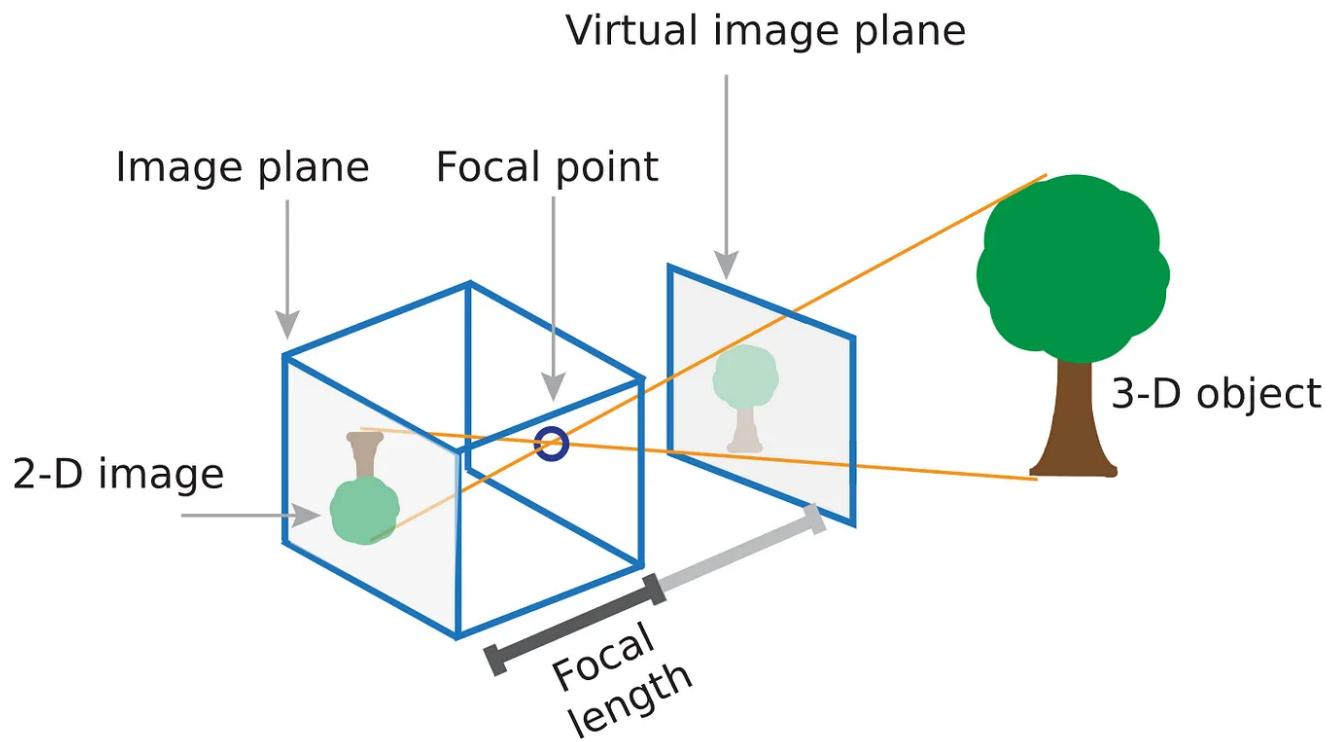
Points expressed in the world frame are projected into the image plane  $[u,v]$  using the perspective projection model  $\boldsymbol{\Pi}$  and the camera intrinsic parameters matrix  $\mathbf{A}$  (also denoted  $\mathbf{K}$  in the literature).

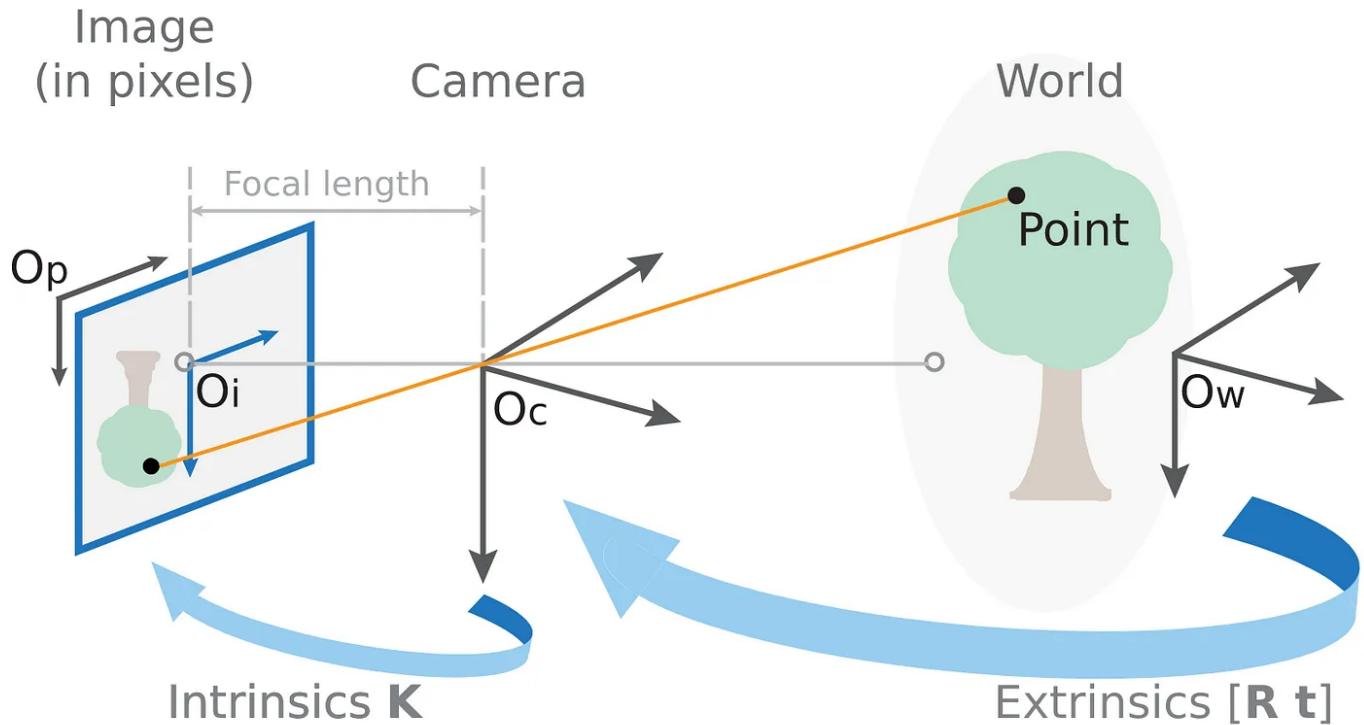
$$\begin{bmatrix} X_c \\ Y_c \\ Z_c \\ 1 \end{bmatrix} = {}^c\mathbf{T}_w \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} X_c \\ Y_c \\ Z_c \\ 1 \end{bmatrix} = \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

Points expressed in the world frame are projected into the camera frame.

The camera intrinsic parameters include the focal length  $f$ , the optical center  $c$ , also known as the principal point, and the skew coefficient (read more [here](#)).





We need an array of 3D object points in the object coordinate space (landmarks of the reference [model in MediaPipe](#)), an array of corresponding image points (2D landmarks of *multiFaceLandmarks* in MediaPipe), intrinsic camera matrix, distortion coefficients (for simplicity we assume that there is no distortion).

Here we use ‘SOLVEPNP\_ITERATIVE’ as a pose computation method which uses a non-linear Levenberg-Marquardt minimization scheme. The initial solution for non-planar “objectPoints” needs at least 6 points and uses the Direct linear transformation algorithm.

The output brings points from the world coordinate system to the camera coordinate system :

- Output rotation vector (see [Rodrigues](#))
- Output translation vector

**Rodrigues' rotation formula** is an efficient algorithm for rotating a vector in space, given an axis and angle of rotation ([Wikipedia](#)).

*cv.Rodrigues*: Converts a rotation matrix to a rotation vector or vice versa.

Here we pass the rotation vector to get a rotation matrix.

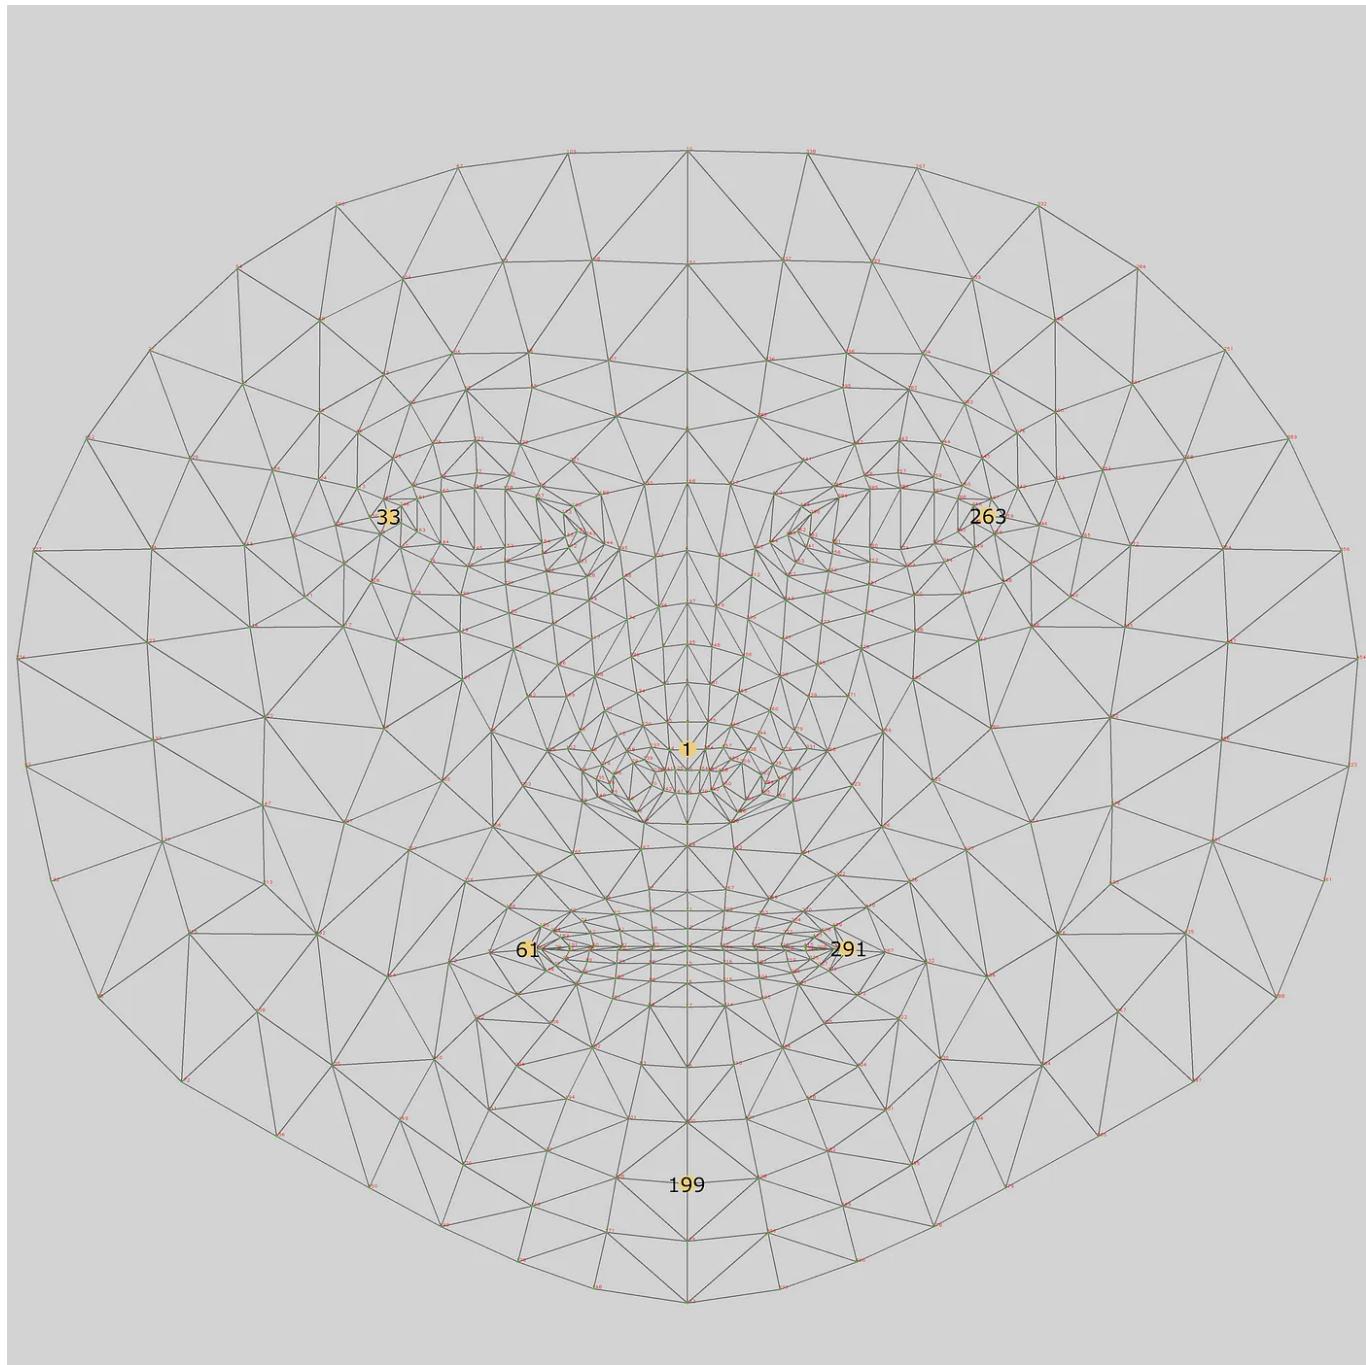
*cv.RQDecomp3x3*: function not found in this OpenCV version, we calculate Euler angles by hand (see [rotationMatrixToEulerAngles](#)).

```
sy = math.sqrt(R[0,0] * R[0,0] + R[1,0] * R[1,0])

singular = sy < 1e-6

if not singular :
    x = math.atan2(R[2,1] , R[2,2])
    y = math.atan2(-R[2,0], sy)
    z = math.atan2(R[1,0], R[0,0])
else :
    x = math.atan2(-R[1,2], R[1,1])
    y = math.atan2(-R[2,0], sy)
    z = 0
```

For each face, MediaPipe FaceMesh contains a bounding box of the detected face and an array of 468 keypoints. Each keypoint or facial landmark has 2D (x, y) or 3D (x, y, z) coordinate locations of facial features, such as lips or eyes corners, points on the eyebrows, irises, and face contours, and intermediate points on the cheeks and forehead.



MediaPipe FaceMesh landmarks with yellow landmarks used for pose estimation (see [here](#))

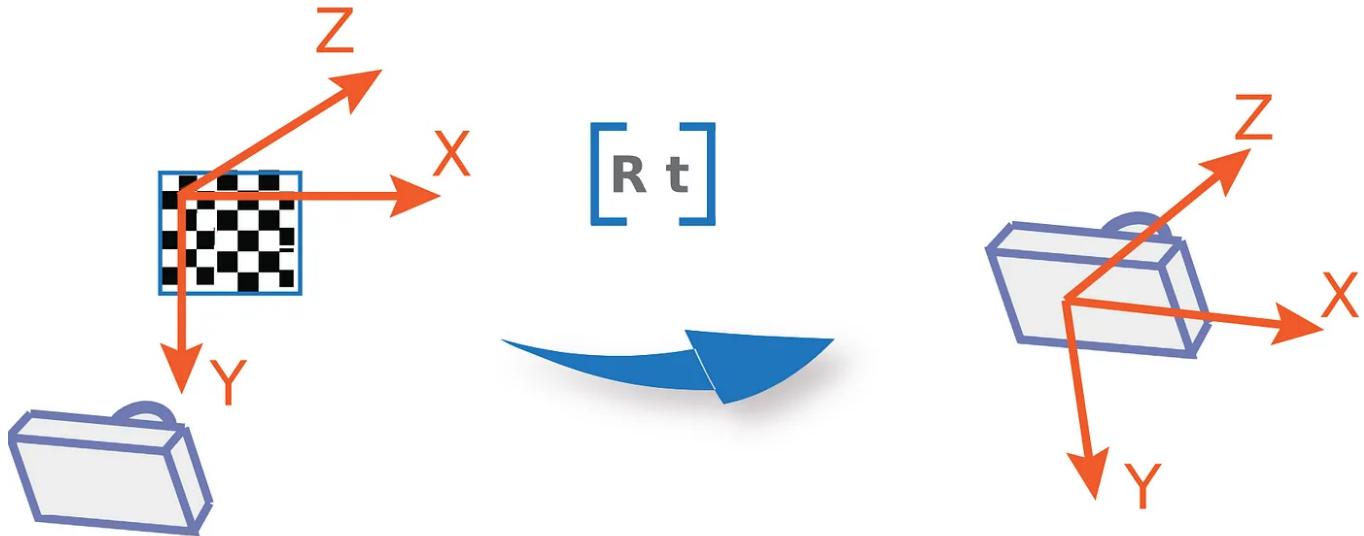
For the keypoints, x and y represent the actual keypoint position in the image pixel space. z represents the depth with the center of the head being the origin, and the smaller the value the closer the keypoint is to the camera. The magnitude of z uses roughly the same scale as x.

The indices of the landmarks we are interested in are 1, 33, 263, 61, 291, and 199 which are evenly distributed on the face.

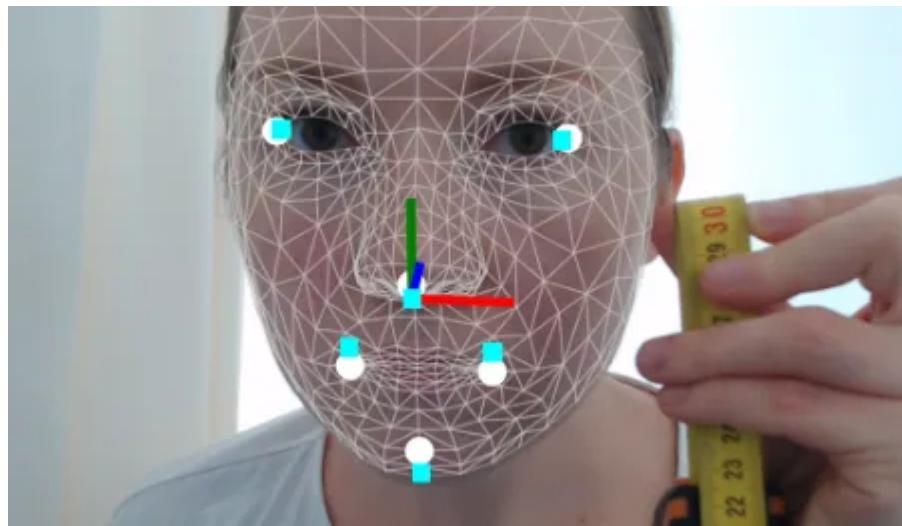
To display the pose of the face as a coordinate system on the nose tip we project the nose landmark to an image plane.

[Open in app](#)[Sign up](#)[Sign In](#)[Search](#) [Write](#)

and translation vector  $t$ , camera matrix  $K$ , and distortion coefficients.



Finally, we draw a line from the 2D nose tip to the projected nose tip.



Final pose estimation result with MediaPipe landmarks in white, the estimated pose as a coordinate system on the nose, and landmarks projected using the estimated pose in cyan.

## References

1. GitHub repository: <https://google.github.io/mediapipe/>
2. Solution “MediaPipe FaceMesh” :  
[https://google.github.io/mediapipe/solutions/face\\_mesh](https://google.github.io/mediapipe/solutions/face_mesh)
3. Nicolai Nielson’s Demo in Python :  
<https://github.com/niconielsen32/ComputerVision/blob/master/headPoseEstimation.py>
4. Mike C. : A head pose estimation Cycle.js demo app using opencv.js and tensorflow.js’ posenet <https://codesandbox.io/s/008olz2wmn?file=/src/index.js:1273-1277>
5. Camera Calibration : <https://fr.mathworks.com/help/vision/ug/camera-calibration.html>
6. Learn OpenCV : <https://learnopencv.com/head-pose-estimation-using-opencv-and-dlib/>
7. Grab the code from here – <https://codepen.io/Susanne-Thierfelder/pen/yLEOQXq>

OpenCV

JavaScript

Computer Vision

Mediapipe

Pose Estimation



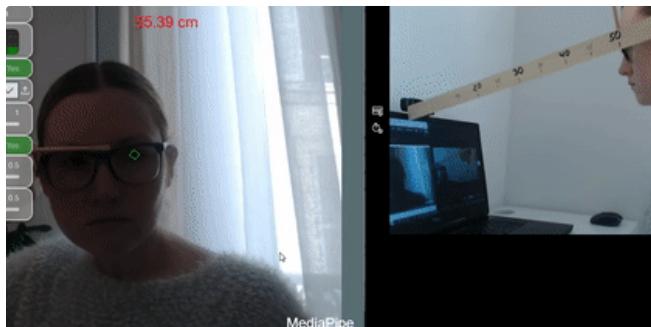
## Written by Susanne Thierfelder

12 Followers

[Follow](#)

---

### More from Susanne Thierfelder

 Susanne Thierfelder

### Create your own depth measuring tool with MediaPipe FaceMesh in...

Getting started with MediaPipe and measuring depth from a single image

4 min read · Nov 3, 2022

 Susanne Thierfelder

### Automate Monthly Wordpress.org Posts for a Food Blog without a...

Use Case: Food Blog using the Blossom Recipes Theme hosted by Bluehost.

2 min read · Oct 5

 15 1 + 3 +

## GREEN FOOD SPOT

NATURALLY SWEET  
TREATS AND  
WHOLESOME  
VEGETARIAN MEALS



Italian focaccia is appreciated for its soft texture, rich flavor, and crispiness. It is often compared to pizza due to its flat shape and topping of herbs and other ingredients, although focaccia is typically thicker than pizza and does not ...

[Continue reading ...](#)[All Recipes](#)[Savory](#)[Sweet](#)[Drinks](#)[Shop](#)

Susanne Thierfelder

## Start a Wordpress Food Blog in 2023

Are you thinking about finally turning your yearlong passion into a blog?

5 min read · Sep 19

 6 2 + 3 +

## Ingredients of the Month September

Susanne Thierfelder

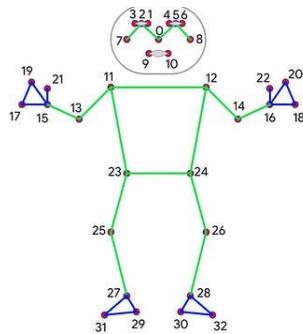
## Automate Monthly Wordpress.org Posts for a Food Blog using a...

This article shows how to automate a Wordpress Blog depending on the current...

3 min read · Sep 27

[See all from Susanne Thierfelder](#)

## Recommended from Medium



- 0. nose
- 1. right eye inner
- 2. right eye
- 3. right eye outer
- 4. left eye inner
- 5. left eye
- 6. left eye outer
- 7. right ear
- 8. left ear
- 9. mouth right
- 10. mouth left
- 11. right shoulder
- 12. left shoulder
- 13. right elbow
- 14. left elbow
- 15. right wrist
- 16. left wrist
- 17. right pinky knuckle #1
- 18. left pinky knuckle #1
- 19. right index knuckle #1
- 20. left index knuckle #1
- 21. right thumb knuckle #2
- 22. left thumb knuckle #2
- 23. right hip
- 24. left hip
- 25. right knee
- 26. left knee
- 27. right ankle
- 28. left ankle
- 29. right heel
- 30. left heel
- 31. right foot index
- 32. left foot index



 MUHAMMD TAYYAB

## Pose Detection Using Mediapipe Solutions (DabMove Detection)...

Introduction

13 min read · May 11

 6 

 +

 Robert John

## How to build a Text to Image App with free AI model in 30 lines of...

Generate amazing Image with AI, a step by step guide to building a Test-to-Image App

6 min read · May 28

 9 

 +

## Lists



### Stories to Help You Grow as a Software Developer

19 stories · 508 saves



### It's never too late or early to start something

15 stories · 190 saves



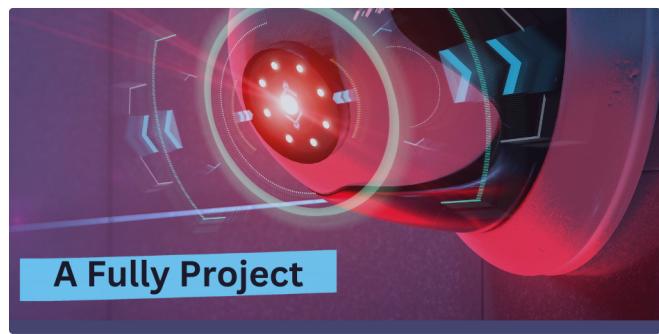
### General Coding Knowledge

20 stories · 523 saves



### Modern Marketing

38 stories · 226 saves





Adithya SM in AI Mind

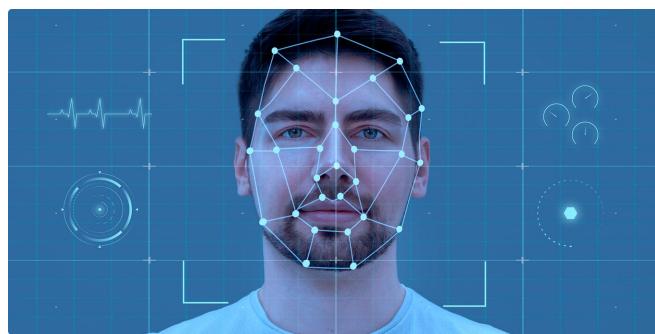


Gal Ben-Evgi

## Real-Time Hand Gesture Recognition Using OpenCV: A...

Introduction

4 min read · Jun 30



KHWAB KALRA

## FACE RECOGNITION

Methods, Implementation, and Practical Examples

5 min read · Jul 11



See more recommendations

## Counting the Number of Objects in an Image: A Machine Learning...

Introduction

5 min read · Sep 1

